

MEASURING POLYGONAL NICENESS*

Sharmila Duppala[†] and David Kraemer[‡]

ABSTRACT. The α -fatness and chord-arc scores as measures of the “niceness” of polygons are developed and explored, with initial theoretical and empirical results.

Contents

1	Introduction	1
2	Theoretical background	2
3	Overview of empirical design	4
4	Empirical results	6
5	Discussion	6
References		8

1 Introduction

Convexity is from many perspectives the “gold standard” of niceness for planar subsets, because it excludes at once large classes of pathological sets altogether. Yet for establishing a quantitative measure of niceness that matches with intuitive or legal understandings, convexity is both too restrictive and too general. The relevant consideration at present revolves around drawing electoral districts. The immense political consequences of favorable (or unfavorable) maps has resulted in the development of creative classes of polygonal regions utilized for this purpose. There exists a strong civic prerogative for providing quantitative measures of niceness. In a concurrent opinion in the 2004 suit *Vieth v. Jubelirer* decided by the U.S. Supreme Court, Justice Anthony Kennedy writes

Because, in the case before us, we have no standard by which to measure the burden appellants claim has been imposed on their representational rights, appellants cannot establish that the alleged political classifications burden those

*This paper constitutes the completion of the final project for AMS 545 “Computational Geometry”.

[†]Department of Computer Science, sduppala@cs.stonybrook.edu

[‡]Department of Applied Mathematics, david.kraemer@stonybrook.edu

same rights. Failing to show that the alleged classifications are unrelated to the aims of apportionment, appellants' evidence at best demonstrates only that the legislature adopted political classifications. That describes no constitutional flaw, at least under the governing Fourteenth Amendment standard. [1]

Ought convexity be considered as such a standard? In practice, convexity is too restrictive (e.g., for coastal or river borders) and too general (e.g., a tiny vertical strip) for this purpose.

The aim of this paper is to propose several alternative measures to determine the niceness of polygonal regions in a quantitatively rigorous way. We introduce two classes of measures, the α -score and the chord- f scores, develop some of their basic properties, and show how they perform on simulated data. The α -fatness quantifies how much a given polygon resembles a ball (with respect to some norm). When f measures the perimeter of a polygon, the chord- f score quantifies the extent to which "local nonconvexity" affects the distribution of the total boundary. Both measures generally reward convex polygons, but punish certain convex "offenders." Of course, there exist nonconvex polygons who score well on both measures as well.

This paper proceeds as follows. Section 2 provides a theoretical overview for the α -fatness and chord- f scores and gives preliminary results on simple examples. Section 3 outlines the algorithms used to perform discretizations and compute the α -fatness and chord- f scores. Section 4 gives the results of numerical simulations for these measures on randomly generated "typical" polygons and provides data about the success of discretization. Section 5 includes a discussion of the main results of the paper and its implications for the electoral districting process.

Source code for this project can be found at <https://github.com/DavidNKraemer/polygonal-niceness>. Documentation, cruft removal, and user interface improvements are presently evolving.

2 Theoretical background

The area of a subset $S \subseteq \mathbb{R}^2$ of the plane is denoted by $\lambda(S)$. We will denote the boundary of S by ∂S . For two points $x, y \in \mathbb{R}^2$, the closed line segment with endpoints in x and y is denoted $[x, y]$. For $p \in [1, \infty]$, we define the p -ball as $B_p(x, \varepsilon) = \{y \in \mathbb{R}^2 : \|x - y\|_p \leq \varepsilon\}$ and we will only be interested in the closed case. Unless otherwise stated, we shall operate in the Euclidean world with $p = 2$. Throughout we shall assume that $P \subseteq \mathbb{R}^2$ is a simple bounded closed polygon. The class of such polygons is given by \mathcal{P} . Its perimeter is denoted by $|\partial P|$.

2.1 The α -fatness score

One approach to measuring the relative "fatness" of a polygon seeks to identify regions of the shape which behave highly unlike balls.

Definition 2.1. The α -fatness score of a polygon P is given by

$$\alpha(P) = \inf \left\{ \frac{\lambda[P \cap B(x, \varepsilon)]}{\lambda[B(x, \varepsilon)]} : \varepsilon > 0, x \in P \right\}$$

where we impose that the ball $B(x, \varepsilon)$ does not contain P .

The constraint that $P \not\subseteq B(x, \varepsilon)$ ensures definiteness: otherwise, let $\varepsilon \rightarrow \infty$ and the ratio $\frac{\lambda[P \cap B(x, \varepsilon)]}{\lambda[B(x, \varepsilon)]}$ always approaches 0. An intuition for $\alpha(P)$ emerges from considering how it behaves on balls.

Proposition 2.2. For any P , $\alpha(P) \leq \frac{1}{4}$, and if $P = B(x_0, r)$ then we have equality.

Proof. To show that $\alpha(P) \leq \frac{1}{4}$, it suffices to show only for convex P , because $P \subseteq (P)$ implies $\lambda[P \cap B(x, \varepsilon)] \leq \lambda[(P) \cap B(x, \varepsilon)]$. Let $d = \text{diam}(P)$ and fix $x, y \in \partial P$ such that $|[x, y]| = d$. Assume without loss of generality that x lies above y , and consider the ball $B(x, d)$. Then P does not meet the upper disk of $B(x, d)$, for otherwise we could choose $z \in P$ in the upper disk of $B(x, d)$ which forms a chord of P longer than d . A similar argument shows that P cannot occupy more than half of the lower disk of $B(x, d)$. Hence

$$\alpha(P) \leq \frac{\lambda[P \cap B(x, d)]}{\lambda[B(x, d)]} \leq \frac{1}{4},$$

as needed.

The fact that $\alpha(B(x_0, r)) = \frac{1}{4}$ relies on a neat result [2] that shows that the L_p ellipsoid \mathcal{E} with radii a, b has area

$$\lambda(\mathcal{E}) = ab \cdot 4 \frac{\Gamma(1 + \frac{1}{p})^2}{\Gamma(1 + \frac{n}{p})}.$$

In our case, $\alpha(B(x_0, r))$ is achieved at a ball on a “corner” point with radius $2r$. □

Hence, the α -score is greater for polygons which exhibit “niceness” characteristics. It is not affected tremendously by local nonconvexity in P so long as P “snugly fits” into a ball shape, which is a global property.

2.2 The chord- f score

A different class of “fatness” measures arises through partitioning P into two (interior disjoint) polygons $P = P' \cup P''$ via chords. A chord is a pair $x, y \in \partial P$ such that the segment $[x, y]$ is contained in the interior of P except at the endpoints x and y . Such a chord defines a partition of $P = P' \cup P''$ by orientation, so that $\partial P'$ is the arc from x to y together with $[x, y]$ and that $\partial P''$ is the arc from y to x together with $[x, y]$.

Definition 2.3. Let $f : \mathcal{P} \rightarrow \mathbb{R}$. For a given $P \in \mathcal{P}$, its *chord- f score* is given by

$$s_f(P) = \inf \{ \max(f(P'), f(P'')) : x, y \in \partial P \}$$

where the chord $[x, y]$ partitions $P = P' \cup P''$.

The intuition for f is some global measure of “cost” associated with respect to P . For example, if $f(P) = |\partial P|$, then s_f identifies the partition (or limiting sequence of partitions) which minimizes the maximum subpolygon perimeter. Indeed, measures such as perimeter or area are the typical choices of f , but any suitable property of P may be employed.

The chord- f score can be understood relatively simply in the context of a minimax game. Max, given a polygon with a chord partitioning it, always chooses the subpolygon associated with a greater f (i.e., the worst cost of the partition). Min, who plays first, tries to find a chord with the least-bad worst cost with respect to f .

The member of this class of scores under this present investigation is $s_{|\partial \cdot|}$, the so-called *chord-arc* score. It is important to note that the length of the chord $[x, y]$ which partitions P is included in the estimation of $s_{|\partial \cdot|}$. One implication of this is that the chord-arc score need not be bounded above by $\frac{|\partial P|}{2}$. We have several facts about $s_{|\partial \cdot|}$ which develop further intuition for the score.

Proposition 2.4. 1. Let R be a rectangle with height h and length ℓ . Assume that $h \leq \ell$.

Then $s_{|\partial \cdot|}(R) = 2h + \ell$.

2. Let P be convex with perimeter $|\partial P| = w$. Then $s_{|\partial \cdot|}(P) \leq \frac{w}{2} + \text{diam}(P)$.

3. Let C be a circle with “perimeter” $2\pi r$. Then $s_{|\partial \cdot|}(C) = (\pi + 2)r$.

Proof. 1. That the optimal partition bisects R is readily apparent, for this minimizes the maximum perimeter length of either sub-rectangle. A case analysis confirms that the preferred strategy for Min is to choose a chord along the smaller dimension (i.e., h) to bisect R , and in this case the maximum arc length is $2h + 2\left(\frac{\ell}{2}\right) = 2h + \ell$, as needed.

2. For P convex, a perimeter-bisecting chord, whose length is at most $\text{diam}(P)$, always exists.

3. This follows from part 2 by noticing that the length of *any* perimeter-bisecting chord is $\text{diam}(C) = 2r$. \square

3 Overview of empirical design

In principal of computing the α and chord- f scores relies on sweeping the continuum of points of P and evaluating the aforementioned measurements. To estimate this process in practice, we employ the following discretizing scheme. Let P be determined by the vertices $[v_1, \dots, v_n]$ given in counterclockwise order. Then, given $\delta > 0$, loop through the vertices of P , checking if $\|v_{i+1} - v_i\| > \delta$. In the case where this holds, we insert $v'_i = \frac{v_i + v_{i+1}}{2}$ after v_i into the polygon and resume the loop. This ensures that for any two adjacent vertices v_i and v_{i+1} in the representation of P we have $\|v_{i+1} - v_i\| \leq \delta$. As $\delta \rightarrow 0$ this estimates ∂P more exactly. All computations are done with respect to the L_∞ norm. The implementation of this procedure is given in Algorithm 1.

Depending on the relative coordinates of the vertex set compared to δ , the worst-case runtime of refinement varies. In the case where the maximum distance is δK , the refinement

115 adds $O(\lg^2 K)$ vertices between the two coordinates. The overall runtime is $O(n \lg^2 K)$ and
 116 the total quantity of vertices after refinement is $O(n + n \lg^2 K)$.

Algorithm 1 The procedure used to refine a given polygon P by a specified $\delta > 0$. This modifies the existing P to satisfy the regularity condition that $\|v_{i+1} - v_i\| \leq \delta$ for each adjacent vertex pair v_{i+1}, v_i on the discretized boundary ∂P .

```

1: procedure REFINEBY( $P, \delta$ )
2:    $P = [v_1, v_2, v_3, \dots, v_n]$ ;
3:    $v = v_1$ 
4:   while  $\text{next}(v) \neq v_1$  do
5:     if  $\|\text{next}(v) - v\|_2 > \delta$  then
6:       Insert  $\text{midpoint}(v, \text{next}(v))$  into  $P$  after  $v$        $\triangleright$  Now  $\text{next}(v)$  is the midpoint

```

Given a δ -refined polygon P , the computation of $\alpha(P)$ is relatively straightforward, with the simplification that points x on which the ratio is computed lie on ∂P . The computation steps through $[v_1, \dots, v_n]$ and for each v_i the radius ε of the minimum-enclosing ∞ -ball centered at v_i is computed. Then for $\varepsilon_k = \frac{k}{100}\varepsilon$, $k = 1, \dots, 100$, the ratio

$$\frac{\lambda[P \cap B_\infty(v_i, \varepsilon_k)]}{B_\infty(v_i, \varepsilon_k)}$$

117 is computed and the smallest such ratio is kept. The minimum over all v_i is given as $\alpha(P)$.
 118 The implementation of the α score is given in Algorithm 2. Though we use a sweep from
 119 $k = 1, \dots, 100$, any suitable range may be selected as an alternative. Unfortunately, there
 120 is no analogous “bisection” technique for determining the minimizing k because this ratio
 121 need not be monotonic. Since P need not be convex, the intersection $P \cap B_\infty(v, \varepsilon_k)$ requires
 122 $O(n \log n)$ steps. For n vertices, the runtime performance is $O(n^2 \log n)$.

Algorithm 2 The α -fatness function.

```

1: function  $\alpha\text{SCORE}(P)$ 
2:    $P = [v_1, v_2, v_3, \dots, v_n]$ ;
3:    $\min = \infty$ 
4:   for  $v \in P$  do
5:      $\varepsilon = \text{radius of minimum enclosing ball of } P \text{ at } v$ 
6:     for  $k = 1, 2, \dots, 100$  do       $\triangleright$  100 is arbitrary
7:        $\varepsilon_k = \frac{k}{100}\varepsilon$ 
8:        $\text{area} = \frac{\lambda[P \cap B_\infty(v, \varepsilon_k)]}{B_\infty(v, \varepsilon_k)}$ 
9:       if  $\min > \text{area}$  then
10:         $\min = \text{area}$ 
11:   return  $\min$ 

```

123 Similarly, given a δ -refined polygon P , the chord- f score is computed by sweeping
 124 all possible chords of P . For each point $v_i, v_j \in \partial P$ with $i < j$, if $[v_i, v_j]$ is indeed a valid
 125 chord of P , the partition $P = P' \cup P''$ is computed and the maximum of $f(P')$ and $f(P'')$
 126 is kept. The minimum such value is given as $s_f(P)$. The implementation of the chord- f
 127 score is given in Algorithm 3. For n vertices, $O(n^2)$ potential chords are computed. To

128 determine if a pair of vertices is a chord takes $O(n)$ time, which leads to an overall runtime
 129 performance of $O(n^3)$.

Algorithm 3 The chord- f function.

```

1: function CHORDSCORE( $P, f$ )
2:    $P = [v_1, v_2, v_3, \dots, v_n]$ ;
3:   for  $v \in P$  do
4:     for  $w \in [next(v), \dots, v_n]$  do
5:       if  $[v, w]$  is a chord of  $P$  then
6:         Partition  $P = P' \cup P''$  by  $[v, w]$ 
7:         submax =  $\max(f(P'), f(P''))$ 
8:         if min > submax then
9:           min = submax
10:  return min

```

130 All of these algorithms are implemented in CGAL [4] with the extensive use of the 2D
 131 Polygon library [3]. For our data generation, we make the minor modification that measured
 132 $s_{|\partial| \cdot |}$ is the ratio of the actual $s_{|\partial| \cdot |}$ score with the perimeter of the overall polygon. This
 133 changes the interpretation of some of the theoretical results from the previous section but
 134 allows for easier interpretation of the data.

135 4 Empirical results

136 We evaluate our measurements on randomly generated polygons with vertices in the unit
 137 square $[0, 1] \times [0, 1]$. We consider collections of such polygons of 10, 50, and 100 vertices.
 138 In Figure 1 and Figure 2, we measure the effect on polygons of 10 vertices of sweeping δ
 139 between 0.05 and 1 on the resulting measurements. Figure 3 shows a comparison of α -fatness
 140 and relative $s_{|\partial| \cdot |\infty|}$ scores for polygons on 10 vertices. Figure 4 shows a comparison of α -
 141 fatness and relative $s_{|\partial| \cdot |1|}$ scores for polygons on 50 vertices. Figure 5 shows a comparison
 142 of α -fatness and relative $s_{|\partial| \cdot |\infty|}$ scores for polygons on 100 vertices.

143 Our theoretical results seem to be confirmed in that “ball”-like polygons get the
 144 highest α -scores, and it seems that polygons with fewer “nonconvexities” score better on
 145 the relative $s_{|\partial| \cdot |\infty|}$ measure. We observe an overall decline in the relative $s_{|\partial| \cdot |}$ scores as
 146 the number of vertices of the polygons increases. This is attributed to the corresponding
 147 increase in the size of the underlying perimeters of each P . The observation is not shared
 148 with the α score, but this can be understood as a sort of “coastline” phenomenon, in which
 149 polygonal perimeters increase without limit while the areas converge.

150 5 Discussion

151 We have developed two classes of measurements for determining the niceness of polygo-
 152 nal regions in the plane and established initial results, both theoretical and empirical, for
 153 evaluating their effectiveness. The α -fatness measures the extent to which the underlying

δ	1	2	3	4	5	6	7	8	9	10
0.05	0.23	0.21	0.11	0.039	0.15	0.038	0.12	0.10	0.12	0.13
0.15	0.23	0.21	0.11	0.039	0.15	0.038	0.12	0.10	0.12	0.13
0.26	0.23	0.21	0.11	0.039	0.15	0.038	0.12	0.10	0.12	0.13
0.36	0.23	0.21	0.11	0.039	0.15	0.038	0.12	0.10	0.12	0.13
0.47	0.23	0.21	0.11	0.039	0.15	0.038	0.12	0.10	0.12	0.13
0.57	0.23	0.21	0.11	0.039	0.15	0.038	0.12	0.10	0.12	0.13
0.68	0.23	0.21	0.11	0.039	0.15	0.038	0.12	0.10	0.12	0.13
0.78	0.23	0.21	0.11	0.039	0.15	0.038	0.12	0.10	0.12	0.13
0.89	0.23	0.21	0.11	0.039	0.15	0.038	0.12	0.10	0.12	0.13
1	0.23	0.21	0.11	0.039	0.15	0.038	0.12	0.10	0.12	0.13

Figure 1: Effect of sweeping δ on measurements of α -score on 10 randomly generated polygons with 10 vertices in $[0, 1]^2$. Unsurprisingly, the refinement of the different polygons has no effect on the α -fatness score, since it introduces no additional area to the polygon and since the “corners” are already present.

154 polygon resembles a ball. It rewards polygons which fill squarely compact regions and pun-
 155 ishes oblong and spread out polygons. The chord-arc score measures the extent to which the
 156 underlying polygon contains significant “local nonconvexities” that are not well distributed
 157 throughout the entire polygon. Interestingly, both measures perform well (even optimally)
 158 on specific types of convex polygons but severely punishes others. We generally find that
 159 these measures conform in some sense to an intuitive understanding of niceness.

160 Extensions on these measures can be examined. The connected α -fatness score, for
 161 example, computes the intersection of $B(x, \varepsilon)$ with the connected component of P containing
 162 x . This might prove a more reliable estimate of intuitive niceness than the conventional α -
 163 fatness score, because it excludes components of P from consideration which are close in a
 164 coarse sense, but not necessarily close in a topological sense. The myriad of possible choices
 165 from which to choose f leaves open a wide space for exploration. A more suitable measure
 166 than simply the perimeter may emerge as a consequence.

167 The next extension of this work is to apply these measures to electoral districts and
 168 evaluate their success at distinguishing between offending and “nice” districts. Given an
 169 increasing corpus of election districts ruled out or struck down by constitutional grounds, it
 170 may be possible to learn a classification model using these measures. Alternatively, it may
 171 be discovered that these measures are sufficient to distinguish between offending and nice
 172 districts in an unsupervised fashion.

173 Whether quantitative measures that conform to intuitive understandings of niceness
 174 is relevant to crafting electoral districts at all remains an open question. General compact-
 175 ness and contiguity are only part of the requirements which an electoral district must satisfy
 176 under the present regime, whereas others, such as those given in the 1965 Voting Rights
 177 Act, bear less consequence from the pure shape of the district. Doubtless these quantita-
 178 tive measures could form part of the evaluation process for determining new districts, but
 179 the actual policy will almost surely involve ensembles of measures, both quantitative and

δ	1	2	3	4	5	6	7	8	9	10
0.05	0.83	0.46	0.83	0.89	0.54	0.67	0.62	1.00	0.86	0.55
0.15	0.79	0.90	0.72	0.60	0.90	0.67	0.99	0.83	0.66	0.90
0.26	1.00	0.75	0.50	0.60	0.90	0.49	0.77	0.52	0.72	0.69
0.36	0.89	1.00	0.50	0.60	0.52	0.49	1.00	0.92	0.74	0.51
0.47	0.61	1.00	0.50	0.75	0.52	0.24	1.00	0.62	1.00	0.51
0.57	0.61	1.00	0.25	0.75	0.52	0.24	1.00	0.62	1.00	0.51
0.68	0.61	1.00	0.25	0.75	0.52	0.24	1.00	0.62	1.00	0.51
0.78	0.61	1.00	0.25	0.75	0.52	0.24	1.00	0.62	1.00	0.51
0.89	0.61	1.00	0.25	0.75	0.52	0.24	1.00	0.62	1.00	0.51
1	0.61	1.00	0.25	0.75	0.52	1.00	1.00	0.62	1.00	0.51

Figure 2: Effect of sweeping δ on measurements of $s_{|\partial \cdot|_\infty}$ on 10 randomly generated polygons with 10 vertices in $[0, 1]^2$. The overall behavior is difficult to summarize. Though one might hope that as $\delta \rightarrow 0$ we see a stabilization in the chord arc score, this does not appear to occur. Moreover, there does not seem to be even a monotonic evolution in the score as $\delta \rightarrow 0$. This is likely due to the fact that smaller values of δ permit stranger nonconvexities to emerge as viable chords.

180 otherwise, in idealized future redistricting.

181 References

- 182 1. Vieth v. Jubelirer 2004.
- 183 2. Wang, X. Volumes of Generalized Unit Balls. *Mathematics Magazine* **78**, 390–395. ISSN:
184 0025570X, 19300980 (2005).
- 185 3. Giezeman, G.-J. & Wesselink, W. in *CGAL User and Reference Manual 4.12* (CGAL
186 Editorial Board, 2018). [https://doc.cgal.org/4.12/Manual/packages.html#](https://doc.cgal.org/4.12/Manual/packages.html#PkgPolygon2Summary)
187 [PkgPolygon2Summary](https://doc.cgal.org/4.12/Manual/packages.html#PkgPolygon2Summary).
- 188 4. The CGAL Project. *CGAL User and Reference Manual 4.12*. [https://doc.cgal.org/](https://doc.cgal.org/4.12/Manual/packages.html)
189 [4.12/Manual/packages.html](https://doc.cgal.org/4.12/Manual/packages.html) (CGAL Editorial Board, 2018).



Figure 3: α scores and $s_{|\partial| \cdot |\infty|}$ scores for randomly generated polygons on 10 vertices.

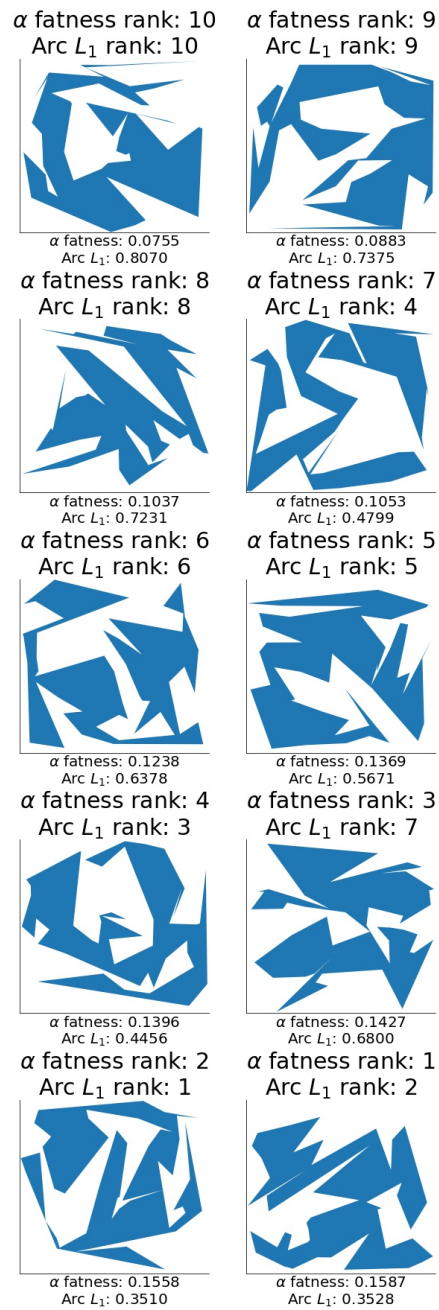


Figure 4: α scores and $s_{|\partial|} \cdot |_1$ scores for randomly generated polygons on 50 vertices.

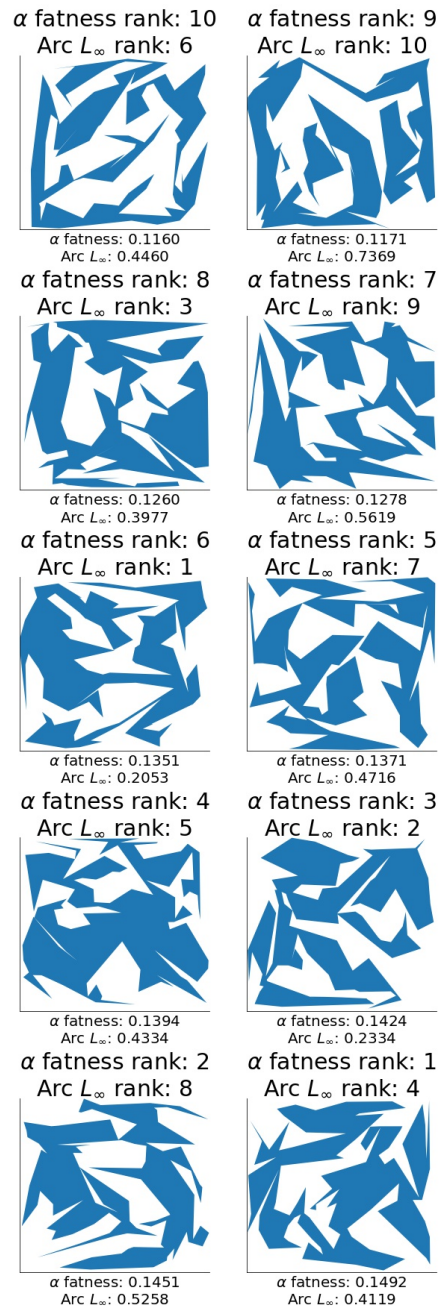


Figure 5: α scores and $s_{|\partial \cdot|_\infty}$ scores for randomly generated polygons on 100 vertices.